

Correction-conditioned media walls for personal taste alignment

Alan N. Pham^{1,*}

¹AO Labs

*Correspondence: aolabs.io

Personal recommendation systems usually learn taste from ratings, clicks, watch time, purchases, or dense behavioral logs. Yum explores a different unit of personalization: the user’s correction. A correction such as “do not show this model, color, pose, ingredient, or source pattern again” is treated as a reusable test that updates a deployed media wall through hard rejects, compact preference memory, source de-duplication, visual-group spacing, and optional model-assisted candidate ranking. The novelty is not another image feed, but a correction-conditioned control layer for personal AI alignment: free-form dissatisfaction is converted into inspectable feed constraints without requiring a large surveillance history. We describe a live static-first website, yum.aolabs.io, that mixes food, cars, and K-pop imagery while preserving category balance and cross-device preference state. Screenshots of the deployed surface are included as artifact evidence. The system is presented as a design study for preference-aware personal media, not as a population-scale recommendation benchmark.

1 Introduction

Recommender systems have long framed personalization as a prediction problem: infer what an individual will rate, click, purchase, or consume next [1, 2, 4, 5]. That framing works when the objective is stable and the signal is abundant. It is weaker when the target is a private aesthetic standard that changes through direct correction: “not this color,” “not this kind of image,” “not this person,” “not this texture,” “not this crop,” or “never cluster several similar items again.”

Yum is built around that second setting. It is a personal media wall whose useful behavior is defined less by aggregate engagement than by whether it stops repeating the user’s rejected mistakes. The core design question is therefore not “how can the system maximize interaction?” but “how can a single user’s corrections become durable operational constraints on a live visual artifact?”

This question matters beyond one website. Human-AI systems increasingly depend on natural-language feedback, preference learning, and alignment from human judgments [8, 9, 10]. However, most deployed personal interfaces still leave users with brittle local settings, opaque ranking models, or short-lived dislike buttons. Yum treats correction as a first-class artifact. A correction is translated into source keys, semantic rejects, group-spacing penalties, category quotas, and preference samples that are small enough to inspect and persistent enough to matter across sessions.

2 Novelty

2.1 From implicit feedback to explicit correction

Implicit-feedback recommenders infer taste from behavior that was not necessarily intended as feedback [4, 5]. A scroll, linger, or click can be ambiguous. Yum instead elevates direct correction. The user may reject an exact item, but the implementation must infer the reusable failure mode: the source may be low-quality, the category may be overrepresented, the subject may be repeated, the image may violate an aesthetic rule, or a cluster may be visually monotonous.

The proposed novelty is *correction-conditioned curation*: a deployed feed architecture in which free-form human correction is compiled into a transparent constraint stack. The stack has four properties.

1. It is **persistent**: a correction can affect a later page load and a different device through remote preference state.
2. It is **inspectable**: hidden keys, hidden samples, kept samples, blocked terms, and group identifiers are ordinary data rather than an opaque behavioral archive.
3. It is **compositional**: hard rejects, category balance, source rotation, model ranking, and layout placement each address a different failure mode.
4. It is **artifact-facing**: the final test is not an offline score but whether the public wall stops showing the rejected pattern.

2.2 Why this is different from a dislike button

A dislike button removes an item. A correction-conditioned system must generalize without overgeneralizing. If the user rejects a blurred K-pop image, the system should prefer higher-resolution images without eliminating the subject. If the user rejects green or purple performance sedans, the system should block those model-color combinations while still allowing clean white compact sedans and adjacent Audi or Mercedes sedans. If several images from one source appear together, the system should treat the cluster as a layout and queueing failure rather than a single bad tile.

This makes the central object of study a feedback-to-artifact loop. The user does not need to become a

60 taxonomy editor. The system must convert correction into durable tests and then verify the deployed
61 surface against those tests.

62 **3 System design**

63 Yum is a static-first site with optional dynamic services. The public wall is served at yum.aolabs.io.
64 The frontend contains known-good seed media, source metadata, category labels, captions, source
65 URLs, shape hints, focus hints, and optional person or car-group identifiers. A small Node service
66 provides preference persistence, model-assisted candidate ranking, and K-pop candidate retrieval. The
67 static wall remains useful if the dynamic service is slow or unavailable.

68 **3.1 Taste state**

69 Preference state is intentionally compact. It contains hidden source keys, hidden samples, kept samples,
70 a version number, and timestamps. Hidden keys remove exact items. Hidden samples provide stronger
71 negative examples. Kept samples provide weak positive context from nearby non-hidden material.
72 This is a privacy-leaning alternative to large behavioral histories: the state is small, inspectable, and
73 biased toward explicit correction rather than passive surveillance.

3.2 Candidate control

Candidate generation and candidate control are separated. Candidate generation finds possible images from static pools and remote sources. Candidate control decides what is allowed to reach the wall. Control includes blocked terms, source de-duplication, category balancing, recent-person spacing, visual-group spacing, car-group spacing, quality thresholds, and optional language-model ranking. This separation is important because visually bad outcomes often arise after a candidate has passed a relevance check. Four similar cars in one view, three portraits of the same person, or repeated low-resolution source patterns are layout-level failures. They must be handled by queue and placement logic, not only by candidate retrieval.

4 Deployed artifact

The deployed wall is shown in Fig. 1 and Fig. 2. The screenshots were captured from the live public URL on May 7, 2026. They show the actual interaction surface: a paper link, the wall itself, mixed subject categories, and a responsive masonry layout.




Figure 1. Desktop view of the deployed Yum wall. The surface opens directly into media rather than a settings dashboard. The visible mix shows the core design tension: category balance, source quality, and visual-group spacing must all hold at once.




Figure 2. Narrow viewport view of the deployed wall. The same feed contract must survive a different column count, where clustering is more noticeable and lazy-loading failures are easier to see.

5 Relationship to prior work

Yum builds on recommender-system work but shifts the evaluation target. Classic surveys and hybrid recommender methods emphasize prediction, filtering, and combining sources of evidence [2, 3]. Implicit-feedback methods formalize preference from observed behavior [4, 5]. Visual recommendation work extends this to appearance-sensitive domains [6]. Human-centered recommender evaluation argues that perceived quality, variety, effort, and privacy affect system experience, not just predictive accuracy [7].

The correction-conditioned wall fits between these traditions and human-AI interaction. Human preference learning and instruction-following models demonstrate that human judgments can train or steer model behavior [8, 9]. Human-AI guidelines emphasize intelligibility, controllability, and graceful handling of user feedback [10]. Vision-language models show that natural language can name visual concepts and support flexible image understanding [11]. Yum applies these ideas at the product-surface scale: the user provides corrections in ordinary language, and the live artifact is updated through explicit feed constraints.

6 Implications

The strongest implication is that small personal systems can be aligned through correction memory without becoming black-box surveillance systems. A personal media wall does not need to store every scroll event to improve. It can store what the user explicitly rejected, what remained nearby, and how those examples map to transparent constraints. This produces a more repairable system: when the wall is wrong, the correction can become a test.

The second implication is that personalization should be evaluated at the artifact boundary. A recommender may rank candidates correctly in isolation while still producing a bad page because similar items cluster, a category disappears, or a visual pattern repeats. Yum therefore treats the rendered wall as the unit of verification.

7 Limitations

This manuscript reports a deployed design study, not a controlled user study. The current evidence is artifact-level: source code, deployed behavior, screenshots, and live verification. No population-scale preference model, benchmark dataset, or statistical generalization claim is made. The system also inherits the limitations of third-party image sources, including uneven metadata quality, unavailable images, source licensing constraints, and inconsistent visual resolution.

The next useful step is not to add more opaque intelligence. It is to measure correction-to-fix latency: after a user gives a rejection pattern, how quickly does the live wall stop violating it, and how often does the generalization accidentally suppress material the user would have liked?

8 Conclusion

Yum argues for correction-conditioned curation as a practical unit of personal AI alignment. A user should not need to repeatedly reject the same failure mode. The system should convert correction into persistent, inspectable constraints and verify the live artifact against those constraints. For personal visual media, this may be a more useful target than engagement maximization: the wall becomes better when it remembers what the user meant, not merely what the user clicked.

References

References

- [1] Resnick, P. & Varian, H. R. Recommender systems. *Communications of the ACM* **40**, 56–58 (1997). doi:10.1145/245108.245121.
- [2] Adomavicius, G. & Tuzhilin, A. Toward the next generation of recommender systems: a survey of the state-of-the-art and possible extensions. *IEEE Transactions on Knowledge and Data Engineering* **17**, 734–749 (2005). doi:10.1109/TKDE.2005.99.

- [3] Burke, R. Hybrid recommender systems: survey and experiments. *User Modeling and User-Adapted Interaction* **12**, 331–370 (2002). doi:10.1023/A:1021240730564.
- [4] Hu, Y., Koren, Y. & Volinsky, C. Collaborative filtering for implicit feedback datasets. In *2008 Eighth IEEE International Conference on Data Mining*, 263–272 (IEEE, 2008). doi:10.1109/ICDM.2008.22.
- [5] Rendle, S., Freudenthaler, C., Gantner, Z. & Schmidt-Thieme, L. BPR: Bayesian personalized ranking from implicit feedback. In *Proceedings of the Twenty-Fifth Conference on Uncertainty in Artificial Intelligence*, 452–461 (AUAI Press, 2009).
- [6] He, R. & McAuley, J. VBPR: visual Bayesian personalized ranking from implicit feedback. In *Proceedings of the Thirtieth AAAI Conference on Artificial Intelligence*, 144–150 (AAAI Press, 2016).
- [7] Knijnenburg, B. P., Willemsen, M. C., Gantner, Z., Soncu, H. & Newell, C. Explaining the user experience of recommender systems. *User Modeling and User-Adapted Interaction* **22**, 441–504 (2012). doi:10.1007/s11257-011-9118-4.
- [8] Christiano, P. F. et al. Deep reinforcement learning from human preferences. *Advances in Neural Information Processing Systems* **30** (2017).
- [9] Ouyang, L. et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems* **35**, 27730–27744 (2022).
- [10] Amershi, S. et al. Guidelines for human-AI interaction. In *Proceedings of the 2019 CHI Conference on Human Factors in Computing Systems*, 1–13 (ACM, 2019). doi:10.1145/3290605.3300233.
- [11] Radford, A. et al. Learning transferable visual models from natural language supervision. In *Proceedings of the 38th International Conference on Machine Learning*, 8748–8763 (PMLR, 2021).
- [12] Heer, J. & Shneiderman, B. Interactive dynamics for visual analysis. *Communications of the ACM* **55**, 45–54 (2012). doi:10.1145/2133806.2133821.

Methods

158

Artifact capture

159

Desktop and narrow-viewport screenshots were captured from the deployed public URL on May 7, 2026 using Playwright. The screenshots are stored in the public asset bundle as assets/yum-wall-desktop-20260507.png and assets/yum-wall-mobile-20260507.png.

Static surface and media metadata

163

The public surface is generated from a static bundle containing the HTML shell, CSS, JavaScript wall renderer, icon assets, preview images, paper files, and figure assets. Media candidates are represented as structured entries with image source, source URL, caption, category, tile shape, focus hint, and optional person or visual-group metadata.

Preference storage

168

Preference state is stored as a compact object containing hidden keys, hidden samples, kept samples, a version number, and update metadata. The browser keeps a local mirror for immediate interaction. The backend stores the same preference shape so state can be restored on another device when the remote service is reachable.

Candidate ranking and layout

173

Candidate ranking is static-first. Known-good seed media is available before any remote service responds. Optional curation requests can pass candidate metadata and preference samples to the configured model. Blocked terms, source keys, category balance, visual groups, person groups, car groups, quality thresholds, and recent placement history are then applied before display. The renderer maintains a background buffer and a fail-open refill path: exact category balance is used when all queues are available, while queued safe media may continue rendering when one category is temporarily refilling so that the wall does not terminate into blank space.

181 Verification

182 The manuscript PDF was generated from this LaTeX source using `pdflatex`. The deployed site
183 was verified by checking that the live paper URL returned a PDF and that the live JavaScript bundle
184 contained the current correction-conditioned curation rules.

185 Data availability

186 No new research dataset was generated. The media metadata required to inspect the public wall is
187 included in the deployed JavaScript bundle and in the Yum application source. Third-party media
188 remains subject to the licensing and availability terms of the original sources linked from the wall.

189 Code availability

190 The application code and this LaTeX source are available in the Yum application repository at
191 github.com/nalalalan/yum-app. The deployed LaTeX source is served at [yum.aolabs.io/yum-preference-](https://yum.aolabs.io/yum-preference-shaped-media-wall.tex)
192 [shaped-media-wall.tex](https://yum.aolabs.io/yum-preference-shaped-media-wall.tex).

193 Acknowledgements

194 No external acknowledgements are declared.

195 Funding statement

196 No external funding is declared.

197 Author contributions

198 A.N.P. conceived the system, implemented the wall and backend, designed the curation and preference-
199 memory behavior, prepared the manuscript, and maintains the public deployment.

Competing interests

200

A.N.P. created and operates AO Labs and Yum.

201

Additional information

202

Supplementary information is not included. Correspondence and requests for materials should be directed to Alan N. Pham through aolabs.io.

203

204